



KubeCon



CloudNativeCon

North America 2017

Building an Edge Computing Platform for Network Services Using Cloud Native Technologies

Stephen Wong, Software Architect, *FutureWei Technologies, Inc.*

Table of Content

- Content
 - Content
 - Content

NFV in the Context of Cloud Native Computing

- Who are we?
 - OPNFV contributors, looking at problem domain of cloud native computing's relevancy in NFV use cases
- NFV — VNF
 - High-level definition of NFV: idea that majority of network functions could be provided by software running on top of COTS servers
 - High-level definition of VNF (virtual network function) — as the term “VNF” will be used throughout the presentation
- The problem we are looking at: can we build VNFs as cloud native applications
 - Scalability
 - Resiliency / Fault Tolerance
 - Composability and reusability
 - Ease of development and deployment

Edge Network Services Applicability

- Edge use cases are suitable for cloud native VNFs due to:
 - 1.Resource constraints: therefore micro-service and containerization provide more optimal resource utilizations
 - 2.High Cost of Edge Maintenance: resiliency and fault tolerant nature of cloud native applications would be beneficial
 - 3.Ease of Deployment: suitable for remote push of new images of components, allows rollback if new component fails locally
 - 4.Replicable: entire deployment can be replicated to different sites

Suitable VNF Type for Cloud Native Application

- Not all VNFs are suitable for being implemented as cloud native applications
- Web backend application is transactional and event-driven by nature, VNFs that follow the same pattern would be natural fit
 - Network functions such as firewall
 - A subset of these network functions is of class transparent proxy — need a unified way to direct traffic to VNF entrance point. This is known as the “service insertion” problem
 - Loadable / composable data path


VNFFG (VNF Forwarding Graph)

- An important NFV orchestration feature is VNFFG
- Cloud native VNF forwarding graph should provide more programmability
- As cloud native application, VNF components communicate with each other via REST APIs / gRPC type mechanism; however, how VNFs are chained is user defined
 - Should everything be API driven?

Terminology

1. CNF: **C**loud-native **N**etwork **F**unction — defined by *operator*, made up of a list of *microservices*
2. Client: initiates traffic which would pass through the edge node
3. Event-handler: a sidecar on a pod. Event-handler is per- μ srv (see definition below), can dynamically load *operator* defined handlers, and implements the SDK stubs
4. μ srv (microservice): a unit of compute service which exposes a set of APIs, the “lego pieces” that would make up a CNF
5. Operator: defines and deploys CNFs

Architecture of Our Approach



Technical Descriptions

I. Event-driven Model

- 1) Event handler running as sidecar in the same pod as microservices
- 2) Microservices expose a set of events, which operators can define and dynamically loaded
- 3) An SDK would be provided by the framework. Operators invoke SDK calls in their custom event handlers, then internally those SDK calls would become gRPC / REST API calls

II. Programmable Datapath

- 1) Takes network policy rules and program data path engine accordingly
- 2) Rule format : { <ip> port : {redirect | copy} => [list of labels]}
- 3) Utilizes eBPF (IOVisor BCC) to load BPF code

Demo Description

1. Operator uses the portal to define a CNF (gateway) which includes multiple μ srvs (policy-mgr, content-inspect), each μ srv publishes a set of events, and operator can choose to implement these event-handlers and customized them via a set of Python SDK methods, and deploy the CNF to a target edge node
2. Previously non running μ srvs (policy-mgr, content-inspect) now running on edge node, with the new event-handler (per μ srv) built/pushed running alongside the core μ srv (running as a container)
3. Client sends an HTTP request — the transport traffic is being redirected to the specified pod associated with the datapath policy's 'redirect' action's label; and subsequently the call sequence matches what the operator specified
4. Operator updates portal to add a μ srv (anti-virus) with updating a previous SDK call (virus_scan SDK call) from an existing event handler (scan_done event from content-inspect)
5. After CNF (re)deploy, the call flow now goes to new μ srv (anti-virus) with the correct RPC call
6. Operator wants to narrow what type of traffic should get virus scan (as opposed to everything), thus updating the event_handlers (remove virus_scan from content-inspect's scan_done, add it as a condition under policy_check from policy-mgr)

Demo

- [To-Be Added: demo diagram]

Roadmap (1): Data Collection

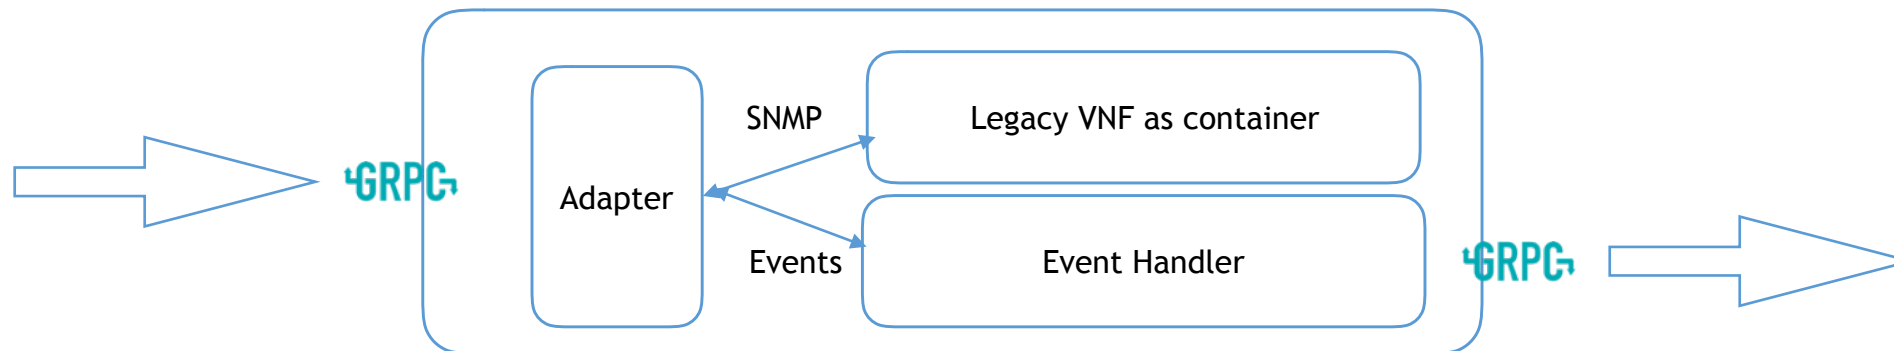
- For edge use cases, it is important to support rich set of data to be collected, also, allowing software-defined filtering of the data being sent to centralized management
- As networking solution, network tracing is important
 - An important reason why eBPF is chosen for datapath engine
 - Particularly IOVisor BCC comes with a rich set of kernel tracing code
- Data correlation: everything throughout the stack (OpenTracing calls, eBPF map, redis objects...etc) all have exact same key (the **cookie** parameter being passed throughout the events / gRPC calls), thus allowing us to bundle data from various sources to be related

Roadmap (2): Multiple Event Handlers per MicroService

- A major drawback of the demo implementation is that the μ srvs are non-sharable, i.e. how the event_handler is invoked per pod makes each μ srv only usable by one CNF
- Moving forward, to allow μ srv to be used by multiple CNFs, cookie filtering value can be used, and event channel can be pub/sub
- Request side needs to make a decision on which CNF a session belongs to (i.e., and most commonly, by client), such info needs to be propagated to **ALL** event_handler on all associated μ srv pods (i.e., like DPM, but for programming the event_handler to subscribe to a topic)
- Programmable session => CNF correlation?

Roadmap (3): Legacy VNF Integration

- There exists a rich and widely deployed set VNFs used by various operators — it is naive and irresponsible to believe VNFs needed to be rewritten to fit this cloud native model
- The two requirements for a μ srv in this framework:
 1. Event generation (and synchronous or asynchronous handlers for those events)
 2. gRPC request handling
- The above requirements do assume the VNFs have some degree of programmability
- One solution to incorporate legacy VNF in this framework would be having another sidecar / container per pod as an adapter to communicate with legacy VNF (as μ srv). The adapter would interface with the legacy VNF (ex: SNMP, CLI) whilst implementing a gRPC server with event generating logic



Roadmap (4): Service Mesh

- If scaling up is needed, this framework matches well with the current service mesh architecture (such as Istio)



Information

- Slack channel: TBD
- Project Clover under OPNFV (<https://wiki.opnfv.org/display/CLOV>)
 - Clover intends to investigate the issues and projects associated with building VNFs as cloud native applications
- Email Address: stephen.kf.wong@gmail.com